

# Architecture Decision Document

---

*This document builds collaboratively through step-by-step discovery. Sections are appended as we work through each architectural decision together.*

## Project Context Analysis

### Requirements Overview

#### **Functional Requirements:**

- Identity & access: Entra ID (EasyAuth) sign-in, minimal profile capture, and user-only data access.
- Programme participation: join/leave/pause with single-programme MVP; opt-in only and no penalties.
- Matching: scheduled cadence, random/lightly constrained matching, avoid repeats, odd-participant handling, idempotent cycles.
- Notifications: match emails with first-move prompt and Teams deep link; email is authoritative fallback.
- Privacy & trust: no individual-level reporting; strict deletion rights; transparency page; event logging of system actions only.
- Administration: role assignment, programme lifecycle controls, aggregate programme health only.

#### **Non-Functional Requirements:**

- Security & privacy: TLS, encryption at rest, least-privilege access, no surveillance analytics, GDPR/UK GDPR/CCPA compliance.
- Reliability: predictable matching execution, atomic runs, clear failure behaviour, idempotent retries.
- Performance: basic web performance expectations (sub-3s load), no real-time constraints.
- Accessibility: WCAG 2.1 AA aspiration, keyboard support, readable contrast.
- Observability: structured logs, error monitoring, aggregate-only insights.
- Email delivery: retry logic, delivery logging, within-60-minute window post-match.

#### **Scale & Complexity:**

- Primary domain: SPA + serverless API + scheduled jobs
- Complexity level: Medium (trust/privacy constraints dominate)
- Estimated architectural components: 6-8 (SPA, auth, API, matching worker, data store, notification service, logging/monitoring, CI/CD)

### Technical Constraints & Dependencies

- Azure Static Web Apps + EasyAuth (no custom auth flows in MVP).
- Azure Functions v4 (.NET isolated) with explicit auth checks per handler.
- Azure Table Storage partitioning by ProgrammId; avoid cross-partition scans.
- Idempotent matching runs keyed by ProgrammId + CycleId.
- OpenAPI.NET generated specs aligned to DTOs and contract tests.
- Email via Azure Communication Services; Teams is secondary.
- Single-tenant MVP with future multi-tenant readiness.
- React 18 + Vite 5 SPA; minimal routing; no state management library in MVP.

## Cross-Cutting Concerns Identified

- Trust contract enforcement (anti-surveillance, opt-in only, no individual metrics).
- Authorisation at every boundary with untrusted claims validation.
- Data minimisation, retention, and deletion workflows.
- Idempotency and failure handling for scheduled matching.
- Aggregate-only reporting constraints for admin/owner views.
- UX tone/clarity as a product promise (low-pressure, low-effort).

## Failure Risks & Early Preventive Guards

- **Trust breach perception:** any hint of surveillance or manager visibility kills adoption -> enforce aggregate-only endpoints and guard-suite tests.
- **Cross-programme data leakage:** partition mistakes expose data -> repository-level key enforcement and negative contract tests.
- **Matching integrity failure:** duplicate runs or partial matches undermine credibility -> idempotency keys and atomic cycle processing.
- **Notification reliability gap:** silent delivery failures feel like "ghosting" -> delivery logging and retry strategy.
- **UX drift into "programme" feel:** extra screens or metrics create obligation -> single-screen core and consistent "no pressure" copy.

## Architectural Drivers

- **Trust by design:** privacy posture drives data minimisation and transparency surfaces.
- **Boundary discipline:** auth/authorisation and partition keys must enforce programme isolation on every query and endpoint.
- **Predictable cadence:** scheduling and idempotent matching are central reliability anchors.
- **Minimal surface area:** few routes and minimal UI states reduce risk and effort.
- **Operational clarity:** logging must explain what the system did, not what users did.

## Stakeholder Lens

- **Participants:** need effortless opt-in/out, zero penalty for silence, and visible proof of non-surveillance -> minimal data capture and clear "what we store" surface.
- **Programme Owners:** need hands-off operation and credible aggregate signals -> aggregate metrics only and reliable cycle completion reporting.
- **Executive Sponsors:** need narrative evidence without privacy risk -> anonymised aggregates and story capture outside the system.
- **Security/Privacy:** require lawful basis, minimisation, retention controls, and auditability -> explicit retention workflows and DPIA-ready documentation.
- **Engineering/Ops:** need low-complexity, debuggable flows -> serverless, deterministic jobs, strong contract tests.

## Cross-Functional Trade-offs

- **PM:** wants validation speed and minimal scope -> architecture stays small, low-config, no platform sprawl.

- **Engineering:** wants reliability and debuggability -> deterministic jobs, explicit failure logging, contract-first APIs.
- **UX:** wants calm, low-pressure surfaces -> avoid extra flows, keep copy consistent, no KPI-styled UI.
- **Privacy/Legal:** wants defensible minimisation and transparency -> strict retention, no behavioural analytics, clear data-handling docs.

## Security & Privacy Lens

- **Attacker view:** highest value targets are identity claims, cross-programme access, and match history -> boundary validation and partition isolation are critical.
- **Defender view:** rely on SWA/EasyAuth for identity, but treat claims as untrusted per request -> explicit auth checks and schema validation everywhere.
- **Auditor view:** needs evidence of minimisation, retention, and deletion controls -> document retention policies and ensure logs exclude user behaviour.

## Expert Panel Notes

- **Serverless architect:** keep Functions thin, avoid framework-building, isolate matching logic into a testable service with idempotency guarantees.
- **Data/privacy specialist:** minimise stored attributes (name, email, department, location), enforce retention windows, prefer anonymised aggregates.
- **Test architect:** prioritise contract tests for auth edge cases, cross-programme access, and idempotent matching; guard tests for anti-surveillance promises.

## Component Failure Modes

- **Auth boundary:** missing/forged claims -> requests must fail closed with explicit 401/403.
- **Programme scoping:** incorrect PartitionKey usage -> accidental cross-programme data access.
- **Matching cycle:** re-run creates duplicate matches -> must detect and no-op on duplicates.
- **Notification delivery:** transient email failures -> retry with observable failure logs.
- **Retention/deletion:** stale personal data persists -> explicit delete workflows and retention checks.
- **Aggregate reporting:** small cohort sizes enable re-identification -> enforce minimum-N thresholds.
- **UX trust drift:** any copy/UI implying obligation or monitoring -> treat as defect, block release.

## Clarifications & Defaults (from stakeholder decisions)

- **Trust contract visibility:** "What we store" is a primary surface on join and first-match screens; full detail behind a link.
- **Cadence promise:** best-effort within a consistent window (e.g., first Tuesday by 10:00 local time); internal alerts on delay, calm external messaging.
- **Aggregate threshold:** minimum-N fixed at 5 for MVP.
- **Retention:** match records retained 12 months; user deletion requests honoured promptly; when leaving, stop matching and retain only minimal non-identifying audit data.
- **Notifications:** email is authoritative; Teams is additive; no SMS.
- **Identity key:** Entra Object ID is canonical; email stored for display/contact only.
- **Time zones:** cadence runs per programme time zone (single-programme MVP).
- **Idempotency key:** ProgrammId + CycleDate (programme time zone).

- **Opt-out semantics:** pause for one cycle and leave programme are both supported; pause is a per-cycle flag.
- **No follow-ups:** ignored matches receive no reminders.

## Starter Template Evaluation

### Primary Technology Domain

Monorepo with three components:

- Web SPA (React/Vite)
- Serverless API (Azure Functions .NET isolated)
- Marketing site (11ty)

### Starter Options Considered

#### 1. **Official Vite React TypeScript template** (for the web app)

- Command (official):

```
pnpm create vite
```

Then select React + TypeScript in prompts, or specify `--template react-ts`.

#### 2. **Azure Functions Core Tools init** (for the API)

- Command:

```
func init collaboratte-api --worker-runtime dotnet-isolated
```

#### 3. **Eleventy (11ty) local install** (for the marketing site)

- Commands (official):

```
pnpm init
pnpm install @11ty/eleventy
npx @11ty/eleventy
```

## Selected Starter: Official Vite + Functions Core Tools + 11ty

### Rationale for Selection:

- Minimal and predictable, aligns with the trust-first, low surface area ethos.
- Keeps UI and content decisions in our control.
- Monorepo enables shared theming assets for both app and marketing site.

## Monorepo Structure (proposed)

```
/apps
  /web          # React + Vite SPA
  /api          # Azure Functions .NET isolated
  /marketing    # 11ty site
/packages
  /theme        # shared design tokens + CSS variables + MUI theme export
```

## Workspaces (pnpm)

- Use `pnpm-workspace.yaml` at the repo root to define workspace packages.

## Shared Theming (lightweight tokens)

- `/packages/theme/tokens.css` for CSS variables used by 11ty
- `/packages/theme/muiTheme.ts` exporting the MUI theme built from the same tokens

## Tooling Additions (confirmed)

### MUI (web app):

```
pnpm add @mui/material @emotion/react @emotion/styled
```

### Playwright (E2E):

```
npm init playwright@latest
```

### Storybook (component workshop):

```
npm create storybook@latest
```

## Deployment (Azure Static Web Apps)

- **Two separate SWA resources:**
  - `app.collabolate.co.uk` -> SPA + API (SWA with `app_location + api_location`)
  - `www.collabolate.co.uk` -> marketing site (SWA with `app_location` only)
- **Plan:** Both SWAs on Free tier for MVP (cost-minimising), with upgrade path to Standard if/when needed.
- **Workflow build configuration:**

- Use `app_location`, `api_location`, and `output_location` in the SWA workflow YAML to match each app folder.
  - `api_location` points to the Functions folder for the app SWA.
- **Routing/auth configuration:**
    - Use `staticwebapp.config.json`; `routes.json` is deprecated.
    - Place `staticwebapp.config.json` under the `app_location`, and ensure the build outputs it to the root of `output_location`.

## Architectural Decisions Provided by Starter

### Language & Runtime:

- React + TypeScript via Vite scaffold.
- Azure Functions .NET isolated worker model for the API.
- 11ty for the marketing site.

### Styling Solution:

- MUI + Emotion in the web app.
- Shared CSS variables/tokens package for 11ty and MUI.

### Testing Framework:

- Playwright for E2E.

### Component Workshop:

- Storybook for isolated UI development.

### Code Organisation:

- pnpm workspaces with `/apps` and `/packages` to share theming across SPA and marketing site.

### Development Experience:

- Vite dev server for the SPA.
- 11ty local build via npx.
- Storybook for UI states and docs.

## Cross-Functional Trade-offs (Starter & Monorepo)

- **PM:** wants speed to validate and clean separation between app and marketing -> two SWAs keep messaging and product surfaces distinct.
- **Engineering:** prefers minimal, predictable scaffolding -> official Vite + Functions init reduces surprises.
- **UX:** wants consistent brand across app and marketing -> shared tokens package is essential.
- **Privacy/Legal:** prefers clear boundary between marketing and authenticated app data -> separate SWAs and domains reduce accidental data coupling.

## Dependency Map (Monorepo & Deployment)

- **Shared theme package** feeds both `/apps/web` (MUI theme) and `/apps/marketing` (`tokens.css`).

- **SWA app+API** depends on repo layout alignment (`app_location`, `api_location`, `output_location`).
- **Marketing SWA** depends on 11ty build output location matching its SWA workflow.
- **Playwright** targets `/apps/web` (app domain), not marketing; marketing can use simple link checks.

## Critical Risks & Mitigations (Starter & Deployment)

- **Monorepo complexity creep:** multiple apps can drift -> keep shared tooling minimal and document boundaries.
- **Theme divergence:** 11ty and MUI could diverge visually -> enforce shared tokens as source of truth.
- **SWA config mismatch:** wrong `app_location/output_location` breaks deploys -> standardise build paths per app and document in workflows.
- **Domain confusion:** users mixing `www` and `app` expectations -> keep navigation boundaries explicit and avoid auth surfaces on marketing.

## Comparative Analysis Matrix (Starter & Deployment)

Option	Maintainability	Trust/Privacy Fit	Delivery Speed	Complexity	Score
<b>Chosen:</b> Official Vite + Functions + 11ty, pnpm, 2x SWA	5	5	4	4	<b>18</b>
Single SWA for app + marketing	3	3	4	3	13
MUI-opinionated starter	3	4	4	3	14
npm workspaces instead of pnpm	4	5	4	4	17

## Pre-mortem (What Could Go Wrong)

- Build pipelines drift between apps -> lock shared tooling and document per-app build outputs.
- Theme package becomes stale -> treat tokens as source of truth and update in one place only.
- Marketing deploy breaks due to 11ty output path mismatch -> standardise output dir and wire in SWA config.
- App SWA fails to route auth correctly -> ensure `staticwebapp.config.json` lands in app build output.

## Clarity Checks (Rubber Duck)

- **Simple:** We will create three apps in one repo, share a theme, and deploy two SWAs (app+API, marketing).
- **Detailed:** `/apps/web` (Vite+React+MUI+Storybook), `/apps/api` (Functions .NET isolated), `/apps/marketing` (11ty), `/packages/theme` for shared tokens.
- **Technical:** pnpm workspaces manage dependencies; SWA workflows point at app and api folders; 11ty output path aligns with marketing SWA.

## Red vs Blue (Starter & Deployment)

- **Red Team:** two SWAs and a monorepo increase config surface area -> higher chance of mis-deploys.

- **Blue Team:** standardised app locations and documented build outputs mitigate; CI checks can validate config.
- **Red Team:** shared theme package could create coupling across apps -> harder to evolve independently.
- **Blue Team:** keep tokens minimal and versioned; avoid deep UI dependencies across app and marketing.

**Note:** Project initialisation using these commands should be the first implementation story.

## Core Architectural Decisions

### Data Architecture

**Decision Priority:** Critical (blocks implementation)

**Database Choice:** Azure Table Storage (Azure.Data.Tables SDK, latest verified 12.11.0). **Rationale:** Aligns with MVP constraints, cost, and simplicity.

**Data Modelling Approach (Table per entity):**

- **Memberships:** PartitionKey = ProgrammeId, RowKey = UserId
- **Cycles:** PartitionKey = ProgrammeId, RowKey = CycleDate (ISO, programme TZ)
- **Matches:** PartitionKey = ProgrammeId, RowKey = MatchId (deterministic from ProgrammeId + CycleDate + sorted UserIds)
- **Events (optional MVP):** PartitionKey = ProgrammeId, RowKey = Timestamp + Guid (only if needed)

**Validation Strategy:** Separate schemas per layer.

- Web: Zod validation.
- API: explicit validation at Functions boundaries.
- No shared schema package in MVP.

**Schema Evolution / Migrations:** Additive changes only; no migrations in MVP.

**Caching:** None in MVP (prioritise observability and simplicity).

**Implications / Cascading Decisions:**

- Deterministic MatchId requires stable canonicalisation of participant IDs.
- CycleDate timezone handling must align with programme TZ decision.
- Event table is optional; if deferred, logging must still satisfy audit requirements.

### Authentication & Security

**Decision Priority:** Critical (blocks implementation)

**Authorisation Pattern:** Role-based, enforced server-side per route.

- Roles: Participant, Programme Owner, Admin.
- Executive Sponsor has no separate role in MVP (same visibility as Programme Owner).
- Role enforcement via `staticwebapp.config.json` and API checks.

## Identity & Claims Handling:

- EasyAuth for identity; validate claims on every request, deny-by-default.
- Never trust client-provided identity; decode SWA client principal header as untrusted input.

## Encryption & Secrets:

- Azure Storage encryption at rest by default; rely on it for Table Storage data.
- No Key Vault in MVP; use SWA/Functions app settings (environment variables) and GitHub Secrets for deployment.
- **Upgrade hook:** SecretsProvider abstraction so Key Vault can be added later without refactoring callers.

## Role Handling (All-Free constraint):

- Keep the concept of roles in API logic, but implement minimal role handling:
  - Default everyone to Participant.
  - Programme Owner/Admin via small allowlist (Entra Object IDs) in Table Storage or app settings.
  - No dynamic role management UI in MVP.
- No manager visibility remains non-negotiable.
- **Upgrade hook:** RolesProvider abstraction to support dynamic role management later.

## API Security:

- Explicit allowlist per route.
- Reject missing/malformed claims; return 401/403.
- No custom auth flows.

## Audit / Event Logging:

- Immutable system-event log only (no user behaviour analytics).
- Log: matching run start/complete, notification attempts/failures, admin config changes.
- Include correlation IDs for debugging.

## Implications / Cascading Decisions:

- Role allowlist source must be documented (Table vs app settings) and secured.

## API & Communication Patterns

### Decision Priority: Critical (blocks implementation)

### API Design: RESTful JSON, noun-based routes.

- Keep route surface small and predictable.
- DTOs only at the boundary.

### API Documentation: OpenAPI.NET generated specs as the source of truth.

- Must stay aligned with DTOs and contract tests.
- Use Microsoft.OpenApi (latest verified 3.1.2).

### Error Handling Standard: Problem Details JSON (RFC 9457, obsoletes 7807).

- Same shape everywhere.
- Human-readable **title** and **detail**.
- Correlation ID included for support/debug.

**Rate Limiting:** None in MVP.

- If abuse appears, add a basic per-user throttle later.

### Implications / Cascading Decisions:

- Standardise error envelope and ensure correlation ID is generated per request.
- OpenAPI specs must be regenerated when DTOs change.

## Frontend Architecture

**Decision Priority:** Important (shapes implementation)

**State Management:** React hooks + Context only. No global state library in MVP.

**Routing:** React Router v6 with minimal route guards.

- Guards only for auth-required vs public routes.
- Clean redirect back to intended page after EasyAuth.
- No client-side permission logic beyond auth presence (server remains authoritative).

**Forms:** Native forms + small helpers only. No React Hook Form.

**Performance:** Keep simple; no explicit code-splitting beyond Vite defaults.

### Implications / Cascading Decisions:

- Auth guard UX must avoid loops and preserve return URLs safely.
- Any future state library must justify added complexity.

## Infrastructure & Deployment

**Decision Priority:** Important (shapes implementation)

**CI/CD:** Standard Azure Static Web Apps GitHub Actions workflows.

- One workflow for app+API SWA.
- One workflow for marketing SWA.

### Environment Configuration:

- **.env** for local dev.
- SWA app settings for deploy-time config.
- No Key Vault in MVP; use app settings and GitHub Secrets.
- **Upgrade hook:** SecretsProvider abstraction to enable Key Vault later.

### Monitoring:

- No Application Insights in MVP; basic platform logs only.

- Log errors and matching-run outcomes only.
- **Upgrade hook:** TelemetryProvider abstraction so App Insights can be enabled later with capped logging.
- Minimal alerting if enabled later: matching run failures, notification send failures, auth failures spike, storage errors.

## Scaling:

- Default consumption/serverless only; no pre-provisioned scaling controls for MVP.

## Implications / Cascading Decisions:

- Alert thresholds and routing must be defined (team email/Slack).
- App settings must be documented per SWA (app vs marketing).

# Implementation Patterns & Consistency Rules

## Pattern Categories Defined

**Critical Conflict Points Identified:** Naming, file structure, API formats, date handling, error/loading patterns.

## Naming Patterns

### Database Naming Conventions:

- **Tables:** PascalCase (`Memberships`, `Cycles`, `Matches`, `Events`)
- **Entity properties:** camelCase (`programmeId`, `cycleDate`, `createdAt`)

### API Naming Conventions:

- **Routes:** plural, noun-based (`/programmes`, `/matches`, `/me`)
- **Route params:** `{id}` (`/programmes/{programmeId}`)
- **JSON fields:** camelCase (`matchedUserId`, `cycleDate`)

### Code Naming Conventions:

- **Components:** PascalCase (`MatchCard`)
- **Frontend files:** kebab-case (`match-card.tsx`)
- **Functions/vars:** camelCase (`getProgrammeById`, `userId`)

## Structure Patterns

### Project Organization:

- Components organised **by feature** (`/features/...`)
- Shared utilities in `/lib`

### Tests:

- Co-located `*.test.ts` / `*.test.tsx` alongside source files

## Format Patterns

## API Response Formats:

- **Direct payloads** only (no `{ data, error }` wrapper)
- **Problem Details** is the only standard error wrapper shape

## Data Exchange Formats:

- JSON fields: camelCase
- Dates: ISO 8601 strings with timezone offsets where relevant
  - Cycle date uses programme TZ

## Communication Patterns

### Events (if/when event log used):

- Event rows use programme-scoped PartitionKey
- Event payload fields in camelCase

## Process Patterns

### Error Handling Patterns:

- Global error boundary for unexpected UI crashes
- Per-call handling for API errors (show inline error states)

### Loading State Patterns:

- Local component state only (no shared loading context)

## Pattern Risks

- **Route naming drift:** agents might reintroduce singular routes -> enforce in code review checklist.
- **Date format drift:** ISO 8601 vs epoch -> contract tests must assert ISO 8601.
- **Error shape drift:** problem-details vs ad-hoc error objects -> lint/test for error envelope consistency.
- **File naming drift:** PascalCase vs kebab-case on frontend -> enforce via linting rule.

## First Principles (Why These Patterns Exist)

- **Consistency beats cleverness:** shared naming/format rules prevent invisible integration bugs.
- **Directness over abstraction:** direct payloads and simple routing reduce accidental divergence.
- **Trust is the product:** error handling and date formats must be predictable and boring.

## Failure Modes

- **Mixed casing in routes/fields:** breaks client/server contract silently.
- **Inconsistent date formats:** invalid comparisons and broken sorting.
- **Non-standard error shapes:** UI error handling fails unpredictably.
- **Misplaced tests:** reduced coverage on contract boundaries.

## Consistency Checks

- Routes are plural and noun-based; params always `{id}`.

- JSON fields and entity properties are camelCase everywhere.
- Frontend file names are kebab-case; component names are PascalCase.
- Problem Details is the only error wrapper; success responses are direct payloads.

## Enforcement Guidelines

### All AI Agents MUST:

- Use PascalCase table names and camelCase for entity/JSON fields.
- Use plural, noun-based API routes with `{id}` params.
- Keep frontend files kebab-case and components PascalCase.
- Use ISO 8601 date strings; Problem Details only for errors.

### Pattern Enforcement:

- Add contract tests for ISO 8601 dates and Problem Details error shape.
- Review checklist includes naming/format rules; deviations require explicit approval.

## Pattern Examples (Condensed)

### Good:

- `GET /programmes/{programmeId} -> { programmeId, cadence, cycleDate }`
- `features/matches/match-card.tsx` exports `MatchCard`

### Anti-Patterns:

- `/programme/123`
- `match_card.tsx`
- `{ data: {...} }` wrapper for success
- epoch timestamps

## Project Structure & Boundaries

### Complete Project Directory Structure

```

collabolate/
├── README.md
├── pnpm-workspace.yaml
├── package.json
├── .gitignore
├── .env.example
└── .github/
    └── workflows/
        ├── swa-app.yml      # app+api SWA
        └── swa-marketing.yml # marketing SWA
└── apps/
    └── web/
        ├── package.json
        ├── vite.config.ts
        └── tsconfig.json

```

```
|- staticwebapp.config.json
|- src/
|   |- app.tsx
|   |- main.tsx
|   |- routes/
|   |- features/
|     |- account/          # includes transparency page
|     |- programmes/
|     |- matches/
|     |- admin/
|   |- lib/
|     |- api/
|     |- auth/
|     |- utils/
|   |- components/
|   |- styles/
|- tests/                  # co-located tests under feature folders
|- api/
|   |- collaboratte-api.sln
|   |- src/
|     |- Collaboratte.Api/
|       |- Program.cs
|       |- Functions/
|         |- Programmes/
|         |- Matches/
|         |- Participation/
|         |- Admin/
|       |- Contracts/
|       |- Services/
|       |- Repositories/
|       |- Models/
|       |- Validation/
|       |- Logging/
|     |- Collaboratte.Api.Tests/
|       |- Functions/      # contract tests live here
|       |- Services/
|       |- Repositories/
|- host.json
|- marketing/
|   |- package.json
|   |- .eleventy.js
|   |- src/
|     |- _data/
|     |- _includes/
|     |- pages/
|     |- assets/
|     |- styles/
|   |- dist/                 # 11ty output (SWA output_location)
|- packages/
|   |- theme/
|     |- package.json
|     |- tokens.css
|     |- muiTheme.ts
|- tests/
```

```
└── e2e/
    ├── playwright.config.ts
    └── specs/
```

## Architectural Boundaries

### API Boundaries:

- REST endpoints exposed in `apps/api/src/Collabolatte.Api/Functions/*`
- Auth boundary enforced per Function (EasyAuth claims validation)

### Component Boundaries:

- Frontend features isolated under `apps/web/src/features/*`
- Shared UI in `apps/web/src/components`
- Shared utilities in `apps/web/src/lib`

### Service Boundaries:

- Domain services in `apps/api/src/Collabolatte.Api/Services`
- Data access via `apps/api/src/Collabolatte.Api.Repositories`

### Data Boundaries:

- Table Storage entities under `apps/api/src/Collabolatte.Api/Models`
- Partition rules enforced in repositories

## Requirements to Structure Mapping

### Feature Mapping (from FR categories):

- **Identity & Access:** `apps/web/src/features/account`,  
`apps/api/src/Collabolatte.Api/Functions/Account`
- **Programme Participation:** `apps/web/src/features/programmes`,  
`apps/api/src/Collabolatte.Api/Functions/Participation`
- **Matching:** `apps/web/src/features/matches`,  
`apps/api/src/Collabolatte.Api/Functions/Matches`
- **Notifications:** `apps/api/src/Collabolatte.Api/Services/Notifications`
- **Administration:** `apps/web/src/features/admin`,  
`apps/api/src/Collabolatte.Api/Functions/Admin`
- **Transparency Page:** `apps/web/src/features/account`

### Cross-Cutting Concerns:

- Auth/claims validation: `apps/api/src/Collabolatte.Api/Validation`
- Logging/event records: `apps/api/src/Collabolatte.Api/Logging`

## Integration Points

### Internal Communication:

- Web app calls API endpoints under `/api/*`.
- API reads/writes Table Storage and sends email via ACS.

### **External Integrations:**

- EasyAuth (SWA) for identity.
- Azure Communication Services for email delivery.

### **Data Flow:**

- Client -> API -> Table Storage
- Scheduler -> Matching service -> Notifications

## File Organization Patterns

### **Configuration Files:**

- `staticwebapp.config.json` must be emitted into the web build output root.
- `.env.example` at repo root.

### **Source Organization:**

- Feature-first in web; feature-grouped Functions in API.

### **Test Organization:**

- Co-located tests in web and API.
- Contract tests under `apps/api/src/Collabolate.Api.Tests/Functions/`.
- Playwright E2E in `/tests/e2e`.

### **Asset Organization:**

- Web assets in `apps/web/src/styles`
- Marketing assets in `apps/marketing/src/assets`

## Development Workflow Integration

### **Development Server Structure:**

- `apps/web` runs Vite dev server.
- `apps/api` runs Functions host locally.
- `apps/marketing` runs 11ty build/watch.

### **Build Process Structure:**

- App SWA builds `apps/web` with API path `apps/api`.
- Marketing SWA builds `apps/marketing` output to `dist`.

### **Deployment Structure:**

- `app.collabolate.co.uk` SWA: `app_location=apps/web, api_location=apps/api`.
- `www.collabolate.co.uk` SWA: `app_location=apps/marketing, output_location=dist`.

## Structure Dependencies

- `/packages/theme` is a shared contract: both `apps/web` and `apps/marketing` depend on it for visual consistency.
- `apps/web` depends on `apps/api` endpoint shapes; API changes must update OpenAPI + contract tests.
- `tests/e2e` depends on stable routes and data fixtures; keep environment config consistent across SWA and local.

## Cross-Functional Structure Trade-offs

- **PM:** wants clear ownership boundaries -> feature-first layout keeps scope visible.
- **Engineering:** wants low merge conflicts -> shared `/packages/theme` reduces duplicated styling.
- **UX:** wants brand consistency across app and marketing -> tokens.css + MUI theme in one place.
- **Privacy/Legal:** wants clear separation between marketing and authenticated app -> distinct apps and SWA workflows.

## Structure Risks

- **Theme package drift:** shared tokens not updated across apps -> add a simple version bump/checklist.
- **API feature sprawl:** Functions folders diverge from routes -> enforce mapping in OpenAPI + tests.
- **Marketing build output mismatch:** 11ty output path doesn't match SWA config -> keep `dist/` consistent and documented.

## Architecture Validation Results

### Coherence Validation

**Decision Compatibility:** All core choices are compatible: React/Vite + SWA EasyAuth + Functions (.NET isolated) + Table Storage + pnpm monorepo. The All-Free constraint is reflected in secrets/roles/monitoring choices.

**Pattern Consistency:** Naming, routing, file structure, and error patterns align with the tech stack and API decisions. Problem Details is the only error wrapper; success responses are direct.

**Structure Alignment:** Project structure supports the decisions (feature-first UI, feature-grouped Functions, shared theme package, co-located tests + E2E).

### Requirements Coverage Validation

**Epic/Feature Coverage:** FR categories map cleanly to `apps/web/src/features/*` and `apps/api/src/Collabolate.Api/Functions/*`.

**Functional Requirements Coverage:** All FR categories (identity, participation, matching, notifications, admin, privacy) are supported by current architecture.

**Non-Functional Requirements Coverage:** Security, privacy, and trust constraints are reflected in auth patterns, data minimisation, and logging rules. Availability/performance constraints are consistent with serverless + simple UI.

### Implementation Readiness Validation

**Decision Completeness:** Critical decisions documented (data, auth, API, frontend, infra). Versions captured where relevant.

**Structure Completeness:** Directory tree is specific and complete across web, API, marketing, theme, and tests.

**Pattern Completeness:** Conflict points are addressed with concrete conventions and enforcement rules.

## Gap Analysis Results

### Important Gap (Resolved):

- **Role allowlist source:** Table Storage is the source of truth (auditable, changeable without redeploys).

### Nice-to-Have:

- Correlation ID format defined as GUID for logs and Problem Details.

## Validation Issues Addressed

- Role allowlist source resolved: **Table Storage**.
- Correlation ID format locked: **GUID**.
- Contract tests for ISO 8601 dates and Problem Details are mandatory gates.

## Architecture Completeness Checklist

### Requirements Analysis

- Project context thoroughly analysed
- Scale and complexity assessed
- Technical constraints identified
- Cross-cutting concerns mapped

### Architectural Decisions

- Critical decisions documented with versions
- Technology stack fully specified
- Integration patterns defined
- Performance considerations addressed

### Implementation Patterns

- Naming conventions established
- Structure patterns defined
- Communication patterns specified
- Process patterns documented

### Project Structure

- Complete directory structure defined
- Component boundaries established
- Integration points mapped

- Requirements to structure mapping complete

## Architecture Readiness Assessment

**Overall Status:** READY FOR IMPLEMENTATION

**Confidence Level:** High

### Key Strengths:

- Trust-first architecture with clear boundaries
- Minimal, consistent conventions to prevent agent drift
- Cost-aware deployment plan with upgrade hooks

### Areas for Future Enhancement:

- Optional Key Vault / App Insights enablement
- Role management UI (post-MVP)

## Implementation Handoff

### AI Agent Guidelines:

- Follow all architectural decisions exactly as documented
- Use implementation patterns consistently across all components
- Respect project structure and boundaries
- Refer to this document for all architectural questions

### First Implementation Priority:

- Initialise monorepo structure + scaffolding per starter decisions

## Architecture Completion Summary

### Workflow Completion

**Architecture Decision Workflow:** COMPLETED

**Total Steps Completed:** 8

**Date Completed:** 2026-01-14

**Document Location:** \_bmad-output/planning-artifacts/architecture.md

## Final Architecture Deliverables

### Complete Architecture Document

- All architectural decisions documented with specific versions where required
- Implementation patterns ensuring AI agent consistency
- Complete project structure with all files and directories
- Requirements-to-structure mapping
- Validation confirming coherence and completeness

### Implementation Ready Foundation

- 5 core architectural decision categories finalised
- 14 consistency and enforcement patterns defined
- 5 architectural components specified (web, api, marketing, theme, e2e)
- 8 functional requirement categories fully supported

## AI Agent Implementation Guide

- Technology stack with verified versions
- Consistency rules that prevent implementation conflicts
- Project structure with clear boundaries
- Integration patterns and communication standards

## Implementation Handoff

**For AI Agents:** This architecture document is your complete guide for implementing collabolate. Follow all decisions, patterns, and structures exactly as documented.

**First Implementation Priority:** Initialise the monorepo structure + scaffolding per starter decisions.

## Development Sequence:

1. Initialise project using documented starter template
2. Set up development environment per architecture
3. Implement core architectural foundations
4. Build features following established patterns
5. Maintain consistency with documented rules

## Quality Assurance Checklist

### Architecture Coherence

- All decisions work together without conflicts
- Technology choices are compatible
- Patterns support the architectural decisions
- Structure aligns with all choices

### Requirements Coverage

- All functional requirements are supported
- All non-functional requirements are addressed
- Cross-cutting concerns are handled
- Integration points are defined

### Implementation Readiness

- Decisions are specific and actionable
- Patterns prevent agent conflicts
- Structure is complete and unambiguous
- Examples are provided for clarity

## Project Success Factors

## **Clear Decision Framework**

Every technology choice was made collaboratively with clear rationale, ensuring all stakeholders understand the architectural direction.

## **Consistency Guarantee**

Implementation patterns and rules ensure that multiple AI agents will produce compatible, consistent code that works together seamlessly.

## **Complete Coverage**

All project requirements are architecturally supported, with clear mapping from business needs to technical implementation.

## **Solid Foundation**

The chosen starter template and architectural patterns provide a production-ready foundation following current best practices.

---

**Architecture Status:** READY FOR IMPLEMENTATION 

**Next Phase:** Begin implementation using the architectural decisions and patterns documented herein.

**Document Maintenance:** Update this architecture when major technical decisions are made during implementation.